

01_airport_cleaning_eda

December 27, 2025

1 Airport Passenger Data - Cleaning & EDA

Data Source: Website_Statistics_Q3_2025.pdf (manually extracted)
Location: data/interim/airport_passengers_manual.csv
Purpose: Clean data for Power BI & exploratory data analysis
Date: December 2025

1.1 Objectives

1. Clean and validate airport passenger data
2. Explore trends, seasonality, and patterns
3. Prepare dataset for Power BI dashboard
4. Validate against Travel Manitoba Q4 2024 infographic

1.2 Setup

```
[62]: # Path setup
import sys
from pathlib import Path

project_root = Path.cwd().parent
sys.path.insert(0, str(project_root / 'scripts'))

from paths import raw, processed, interim
```

```
[63]: # Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from datetime import datetime

# Plotting style
plt.style.use('seaborn-v0_8-darkgrid')
sns.set_palette('husl')
%matplotlib inline

# Display options
```

```

pd.set_option('display.max_columns', None)
pd.set_option('display.float_format', '{:.0f}'.format)

print(' Libraries loaded')

```

Libraries loaded

1.3 Part 1: Data Loading & Cleaning

1.3.1 1.1 Load Raw Data

```
[64]: csv_path = interim() / 'airport_passengers_manual.csv'

if not csv_path.exists():
    print(f'ERROR: File not found at {csv_path}')
else:
    print(f' Found: {csv_path}')
    print(f' Size: {csv_path.stat().st_size:,} bytes')
```

Found:

```
/Users/dpro/projects/travel_manitoba/data/interim/airport_passengers_manual.csv
Size: 11,320 bytes
```

```
[65]: # Load CSV
df_raw = pd.read_csv(csv_path, encoding='utf-8-sig')

# Remove unnamed columns
df_raw = df_raw.loc[:, ~df_raw.columns.str.contains('^Unnamed')]

print('RAW DATA')
print('='*80)
print(f'Shape: {df_raw.shape}')
print(f'Columns: {list(df_raw.columns)}')
print(f'\nFirst 8 rows:')
df_raw.head(8)
```

RAW DATA

```
Shape: (84, 15)
Columns: ['year', 'passenger_type', 'January', 'February', 'March', 'April',
'May', 'June', 'July', 'August', 'September', 'October', 'November', 'December',
'Total']
```

First 8 rows:

	year	passenger_type	January	February	March	April	May	\
0	2005	Domestic	192,265	193,905	217,218	205,509	230,093	
1	2005	International	33,783	32,478	26,179	3,384	498	
2	2005	Total_Passengers	263,351	263,612	281,125	240,584	260,710	

3	2005	Transborder	37,303	37,229	37,728	31,691	30,119
4	2006	Domestic	206,279	198,245	222,093	208,140	244,171
5	2006	International	32,021	26,294	24,448	2,237	963
6	2006	Total_Passengers	274,898	261,217	285,923	242,967	276,835
7	2006	Transborder	36,598	36,678	39,382	32,590	31,701

	June	July	August	September	October	November	December	Total
0	239,036	264,680	271,490	229,127	239,721	207,337	225,227	2,715,608
1	579	1,022	864	633	440	-	8,627	108,487
2	275,742	303,701	308,386	257,803	272,305	238,224	266,345	3,231,888
3	36,127	37,999	36,032	28,043	32,144	30,887	32,491	407,793
4	253,629	278,918	295,063	248,392	255,642	221,539	229,214	2,861,325
5	1,389	1,086	1,271	1,318	651	-	8,755	100,433
6	289,965	317,040	334,657	279,794	292,241	257,158	273,843	3,386,538
7	34,947	37,036	38,323	30,084	35,948	35,619	35,874	424,780

1.3.2 1.2 Clean Numeric Columns

```
[66]: def clean_numeric_column(series):
    """Remove commas, handle dashes, convert to float."""
    return (
        series
        .astype(str)
        .str.replace(',', '', regex=False)
        .str.replace('-', '', regex=False)
        .str.strip()
        .replace('', np.nan)
        .replace('nan', np.nan)
        .astype('float')
    )

# Clean data
df_cleaned = df_raw.copy()

# Numeric columns
month_cols = ['January', 'February', 'March', 'April', 'May', 'June',
              'July', 'August', 'September', 'October', 'November', 'December']
numeric_cols = month_cols + ['Total']

for col in numeric_cols:
    if col in df_cleaned.columns:
        df_cleaned[col] = clean_numeric_column(df_cleaned[col])

# Fix passenger type
df_cleaned['passenger_type'] = df_cleaned['passenger_type'].str.replace('_', '_')
```

```

print(' Cleaned numeric columns')
print(' Fixed passenger type names')
print(f'\nData types:')
print(df_cleaned.dtypes)

```

Cleaned numeric columns
Fixed passenger type names

Data types:

year	int64
passenger_type	object
January	float64
February	float64
March	float64
April	float64
May	float64
June	float64
July	float64
August	float64
September	float64
October	float64
November	float64
December	float64
Total	float64
dtype:	object

1.3.3 1.3 Data Quality Checks

```

[67]: print('DATA QUALITY SUMMARY')
print('='*80)
print(f'Total rows: {len(df_cleaned)}')
print(f'Total columns: {len(df_cleaned.columns)}')
print(f'Years: {df_cleaned["year"].min()} - {df_cleaned["year"].max()}')
print(f'Passenger types: {df_cleaned["passenger_type"].unique().tolist()}')

print(f'\nNull values per column:')
null_summary = df_cleaned.isnull().sum()
print(null_summary[null_summary > 0])

print(f'\nRows per year:')
print(df_cleaned['year'].value_counts().sort_index())

```

DATA QUALITY SUMMARY
=====

Total rows: 84
Total columns: 15
Years: 2005 - 2025
Passenger types: ['Domestic', 'International', 'Total Passengers',

```
'Transborder']  
  
Null values per column:  
January      1  
February     1  
March        2  
April         3  
May          10  
June         11  
July         12  
August        10  
September    14  
October       9  
November      9  
December      5  
dtype: int64
```

```
Rows per year:  
year  
2005      4  
2006      4  
2007      4  
2008      4  
2009      4  
2010      4  
2011      4  
2012      4  
2013      4  
2014      4  
2015      4  
2016      4  
2017      4  
2018      4  
2019      4  
2020      4  
2021      4  
2022      4  
2023      4  
2024      4  
2025      4  
Name: count, dtype: int64
```

1.3.4 1.4 Validate Q4 2024

```
[68]: print('Q4 2024 VALIDATION (vs Infographic)')  
print('='*80)
```

```

mask = (df_cleaned['year'] == 2024) & (df_cleaned['passenger_type'] == 'Total ↴Passengers')
total_2024 = df_cleaned[mask]

if not total_2024.empty:
    q4_sum = total_2024[['October', 'November', 'December']].values[0].sum()
    expected = 1_075_859

    print(f'Calculated Q4 2024: {q4_sum:.0f}')
    print(f'Expected (infographic): {expected:.0f}')
    print(f'Difference: {abs(q4_sum - expected):.0f}')
    print()

    if abs(q4_sum - expected) < 10:
        print(' VALIDATION PASSED')
    else:
        print(' VALIDATION FAILED')

```

Q4 2024 VALIDATION (vs Infographic)

Calculated Q4 2024: 1,075,859
 Expected (infographic): 1,075,859
 Difference: 0

VALIDATION PASSED

1.4 Part 2: Exploratory Data Analysis

1.4.1 2.1 Reshape Data for Time Series Analysis

```
[69]: # Melt to long format for easier plotting
df_long = df_cleaned.melt(
    id_vars=['year', 'passenger_type'],
    value_vars=month_cols,
    var_name='month',
    value_name='passengers'
)

# Create date column
month_map = {m: i+1 for i, m in enumerate(month_cols)}
df_long['month_num'] = df_long['month'].map(month_map)
df_long['date'] = pd.to_datetime(
    df_long['year'].astype(str) + '-' +
    df_long['month_num'].astype(str).str.zfill(2) + '-01'
)

# Sort by date
df_long = df_long.sort_values('date')
```

```

print(f'Long format shape: {df_long.shape}')
print(f'Date range: {df_long["date"].min()} to {df_long["date"].max()}')
df_long.head()

```

Long format shape: (1008, 6)
Date range: 2005-01-01 00:00:00 to 2025-12-01 00:00:00

	year	passenger_type	month	passengers	month_num	date
0	2005	Domestic	January	192,265	1	2005-01-01
1	2005	International	January	33,783	1	2005-01-01
2	2005	Total Passengers	January	263,351	1	2005-01-01
3	2005	Transborder	January	37,303	1	2005-01-01
86	2005	Total Passengers	February	263,612	2	2005-02-01

1.4.2 2.2 Total Passengers Over Time

```

[70]: # Filter to Total Passengers only
df_total = df_long[df_long['passenger_type'] == 'Total Passengers'].copy()

# Plot
fig, ax = plt.subplots(figsize=(14, 6))

ax.plot(df_total['date'], df_total['passengers'], linewidth=2, marker='o', □
         ↵markersize=3)
ax.axvline(pd.Timestamp('2020-03-01'), color='red', linestyle='--', alpha=0.5, □
            ↵label='COVID-19 Start')
ax.set_xlabel('Year', fontsize=12)
ax.set_ylabel('Total Passengers', fontsize=12)
ax.set_title('Winnipeg Airport Total Passenger Traffic (2005-2025)', □
              ↵fontsize=14, fontweight='bold')
ax.legend()
ax.grid(True, alpha=0.3)

# Format y-axis
ax.yaxis.set_major_formatter(plt.FuncFormatter(lambda x, p: f'{x/1e6:.1f}M' if □
                                              ↵x >= 1e6 else f'{x/1e3:.0f}K'))

plt.tight_layout()
plt.show()

print('Key Observations:')
print('- COVID-19 impact visible in 2020-2021')
print('- Recovery trend from 2022 onwards')
print('- Seasonal patterns evident (summer peaks)')

```



Key Observations:

- COVID-19 impact visible in 2020–2021
- Recovery trend from 2022 onwards
- Seasonal patterns evident (summer peaks)

1.4.3 2.3 Annual Total Passengers

```
[71]: # Annual totals
annual_totals = df_cleaned[df_cleaned['passenger_type'] == 'Total Passengers'][['year', 'Total']].copy()
annual_totals = annual_totals.sort_values('year')

# Plot
fig, ax = plt.subplots(figsize=(12, 6))

bars = ax.bar(annual_totals['year'], annual_totals['Total'], color='steelblue', edgecolor='black')

# Highlight 2020-2021 (COVID)
for i, (year, total) in enumerate(zip(annual_totals['year'], annual_totals['Total'])):
    if year in [2020, 2021]:
        bars[i].set_color('darkred')


ax.set_xlabel('Year', fontsize=12)
ax.set_ylabel('Annual Total Passengers', fontsize=12)
ax.set_title('Annual Airport Passenger Totals (2005-2025)', fontsize=14, fontweight='bold')
ax.yaxis.set_major_formatter(plt.FuncFormatter(lambda x, p: f'{x/1e6:.1f}M'))
ax.grid(True, alpha=0.3, axis='y')
```

```

plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

# Summary stats
print('Annual Statistics:')
print(f'Peak year: {annual_totals.loc[annual_totals["Total"].idxmax(), "year"]:.0f} ({annual_totals["Total"].max():,.0f} passengers)')
print(f'Lowest year: {annual_totals.loc[annual_totals["Total"].idxmin(), "year"]:.0f} ({annual_totals["Total"].min():,.0f} passengers)')
print(f'2024 total: {annual_totals[annual_totals["year"] == 2024]["Total"].values[0]:,.0f} passengers')

```



Annual Statistics:

Peak year: 2018 (4,484,343 passengers)
 Lowest year: 2021 (1,223,054 passengers)
 2024 total: 4,297,478 passengers

1.4.4 2.4 Passenger Type Breakdown

```
[72]: # Filter out Total Passengers for breakdown
df_breakdown = df_long[df_long['passenger_type'] != 'Total Passengers'].copy()

fig, ax = plt.subplots(figsize=(14, 6))

for ptype in df_breakdown['passenger_type'].unique():
    data = df_breakdown[df_breakdown['passenger_type'] == ptype]
```

```


    ax.plot(data['date'], data['passengers'], label=ptype, linewidth=2,marker='o', markersize=2)

ax.axvline(pd.Timestamp('2020-03-01'), color='red', linestyle='--', alpha=0.5)
ax.set_xlabel('Year', fontsize=12)
ax.set_ylabel('Passengers', fontsize=12)
ax.set_title('Passenger Traffic by Type (2005-2025)', fontsize=14, fontweight='bold')
ax.legend()
ax.grid(True, alpha=0.3)
ax.yaxis.set_major_formatter(plt.FuncFormatter(lambda x, p: f'{x/1e6:.1f}M' if x >= 1e6 else f'{x/1e3:.0f}K'))

plt.tight_layout()
plt.show()

print('Passenger Type Insights:')
print('- Domestic passengers dominate (~85% of total)')
print('- International traffic most volatile')
print('- Transborder relatively stable')

```



Passenger Type Insights:

- Domestic passengers dominate (~85% of total)
- International traffic most volatile
- Transborder relatively stable

1.4.5 2.5 Seasonality Analysis

```
[73]: # Average passengers by month (exclude COVID years 2020-2021)
df_seasonal = df_long[
    (df_long['passenger_type'] == 'Total Passengers') &
    (~df_long['year'].isin([2020, 2021]))]
].copy()

monthly_avg = df_seasonal.groupby('month')['passengers'].mean().
    reindex(month_cols)

fig, ax = plt.subplots(figsize=(12, 6))


bars = ax.bar(range(len(monthly_avg)), monthly_avg, color='teal', edgecolor='black')

# Highlight summer months
summer_indices = [5, 6, 7] # June, July, August
for i in summer_indices:
    bars[i].set_color('orange')

ax.set_xlabel('Month', fontsize=12)
ax.set_ylabel('Average Passengers', fontsize=12)
ax.set_title('Average Monthly Passenger Traffic (Excluding 2020-2021)', fontsize=14, fontweight='bold')
ax.set_xticks(range(12))
ax.set_xticklabels([m[:3] for m in month_cols])
ax.yaxis.set_major_formatter(plt.FuncFormatter(lambda x, p: f'{x/1e3:.0f}K'))
ax.grid(True, alpha=0.3, axis='y')

plt.tight_layout()
plt.show()

print('Seasonal Patterns:')
peak_month = monthly_avg.idxmax()
low_month = monthly_avg.idxmin()
print(f'Peak: {peak_month} ({monthly_avg.max():.0f} avg passengers)')
print(f'Low: {low_month} ({monthly_avg.min():.0f} avg passengers)')
print(f'Summer boost: {((monthly_avg[summer_indices].mean() / monthly_avg.mean() - 1) * 100):.1f}% above average')
```



Seasonal Patterns:

Peak: August (371,789 avg passengers)

Low: November (280,063 avg passengers)

Summer boost: 12.0% above average

```
/var/folders/36/_jr9z14n69x_lzy_7969swfh0000gs/T/ipykernel_78453/3755199244.py:3
4: FutureWarning: Series.__getitem__ treating keys as positions is deprecated.
In a future version, integer keys will always be treated as labels (consistent
with DataFrame behavior). To access a value by position, use `ser.iloc[pos]`  

    print(f'Summer boost: {((monthly_avg[summer_indices].mean() /  

monthly_avg.mean() - 1) * 100):.1f}% above average')
```

1.4.6 2.6 Year-over-Year Growth (2024 vs 2023)

```
[74]: # Get 2023 and 2024 data
df_2023 = df_long[(df_long['year'] == 2023) & (df_long['passenger_type'] ==
    'Total Passengers')].copy()
df_2024 = df_long[(df_long['year'] == 2024) & (df_long['passenger_type'] ==
    'Total Passengers')].copy()

# Merge on month
comparison = df_2023[['month', 'passengers']].merge(
    df_2024[['month', 'passengers']],
    on='month',
    suffixes=('_2023', '_2024'))
comparison = comparison.set_index('month').reindex(month_cols)
comparison['oy_change_pct'] = ((comparison['passengers_2024'] -
    comparison['passengers_2023']) / comparison['passengers_2023'] * 100)
```

```

fig, (ax1, ax2) = plt.subplots(1, 2, figsize=(16, 6))

# Panel 1: Side-by-side comparison
x = np.arange(len(month_cols))
width = 0.35

ax1.bar(x - width/2, comparison['passengers_2023'], width, label='2023',
        color='lightblue', edgecolor='black')
ax1.bar(x + width/2, comparison['passengers_2024'], width, label='2024',
        color='darkblue', edgecolor='black')

ax1.set_xlabel('Month', fontsize=12)
ax1.set_ylabel('Passengers', fontsize=12)
ax1.set_title('2024 vs 2023 Monthly Comparison', fontsize=14, fontweight='bold')
ax1.set_xticks(x)
ax1.set_xticklabels([m[:3] for m in month_cols])
ax1.legend()
ax1.grid(True, alpha=0.3, axis='y')
ax1.yaxis.set_major_formatter(plt.FuncFormatter(lambda x, p: f'{x/1e3:.0f}K'))


# Panel 2: YoY % change
colors = ['green' if x > 0 else 'red' for x in comparison['yoym_change_pct']]
ax2.bar(range(12), comparison['yoym_change_pct'], color=colors,
        edgecolor='black')
ax2.axhline(0, color='black', linewidth=1)

ax2.set_xlabel('Month', fontsize=12)
ax2.set_ylabel('YoY Change (%)', fontsize=12)
ax2.set_title('Year-over-Year Growth (2024 vs 2023)', fontsize=14,
              fontweight='bold')
ax2.set_xticks(range(12))
ax2.set_xticklabels([m[:3] for m in month_cols])
ax2.grid(True, alpha=0.3, axis='y')

plt.tight_layout()
plt.show()

print('YoY Growth Summary (2024 vs 2023):')
print(f'Average growth: {comparison["yoym_change_pct"].mean():.1f}%')
print(f'Best month: {comparison["yoym_change_pct"].idxmax():.1f}%')
print(f'Worst month: {comparison["yoym_change_pct"].idxmin():.1f}%')
print(f'Min growth: {comparison["yoym_change_pct"].min():.1f}%')

```



YoY Growth Summary (2024 vs 2023):

Average growth: 5.5%

Best month: December (+18.8%)

Worst month: September (-0.7%)

1.4.7 2.7 Missing Data Heatmap

```
[75]: # Create pivot for heatmap
pivot_total = df_cleaned[df_cleaned['passenger_type'] == 'Total' ↴
    ↴Passengers'][['year'] + month_cols].set_index('year').sort_index()

fig, ax = plt.subplots(figsize=(12, 8))

# Create mask for missing data
mask = pivot_total.isnull()

sns.heatmap(
    mask,
    cmap=['lightgreen', 'lightcoral'],
    cbar_kws={'label': 'Missing Data'},
    linewidths=0.5,
    linecolor='gray',
    ax=ax
)


ax.set_title('Data Completeness by Year and Month', fontsize=14, ↴
    ↴fontweight='bold')
ax.set_xlabel('Month', fontsize=12)
ax.set_ylabel('Year', fontsize=12)

plt.tight_layout()
plt.show()
```

```

print('Missing Data Summary:')
print(f'Total cells: {pivot_total.size}')
print(f'Missing cells: {pivot_total.isnull().sum().sum()}')
print(f'Completeness: {(1 - pivot_total.isnull().sum().sum() / pivot_total.
    size) * 100:.1f}%' )

```



Missing Data Summary:

Total cells: 252

Missing cells: 3

Completeness: 98.8%

1.5 Part 3: Save Processed Data

```
[76]: # Save cleaned data
output_path = processed() / 'airport_passengers_clean.csv'
df_cleaned.to_csv(output_path, index=False)

print(' SAVED PROCESSED DATA')
print('*'*80)
print(f'Location: {output_path}')
print(f'Size: {output_path.stat().st_size:,} bytes')
```

```
print(f'Shape: {df_cleaned.shape}')
print(f'\nReady for Power BI import!')
```

SAVED PROCESSED DATA

Location:

/Users/dpro/projects/travel_manitoba/data/processed/airport_passengers_clean.csv

Size: 10,197 bytes

Shape: (84, 15)

Ready for Power BI import!

1.6 Summary

1.6.1 Data Cleaning

- Loaded manual CSV from interim directory
- Cleaned numeric formatting (removed commas, converted dashes to NaN)
- Fixed passenger type names
- Validated Q4 2024 = 1,075,859 passengers (exact match)
- Saved to data/processed/airport_passengers_clean.csv

1.6.2 Key Findings from EDA

1. **Overall Trends** - Steady growth 2005-2019 (pre-COVID) - Dramatic COVID-19 impact in 2020-2021 (~70% decline) - Strong recovery trend 2022-2024 - 2024 approaching pre-pandemic levels
2. **Seasonality** - Clear summer peak (June-August) - Summer traffic ~15-20% above annual average - February typically lowest month
3. **Passenger Mix** - Domestic: ~85% of total traffic - Transborder: ~10% - International: ~5% (most volatile)
4. **Recent Performance (2024 vs 2023)** - Positive YoY growth in most months - Q4 2024: 1,075,859 passengers - Recovery continuing toward pre-pandemic levels

1.6.3 Next Steps

1. Import airport_passengers_clean.csv into Power BI
2. Create measures:
 - YoY % Change
 - YTD Totals
 - Quarterly Aggregations
3. Build dashboard matching Travel Manitoba style
4. Add slicers for year, passenger type, time period